

Response to Reviewers — Paper 30

“Just Ask for a Table: A Thirty-Token User Prompt

Defeats Sponsored Recommendations in Twelve LLMs”

We thank both reviewers for their careful and constructive reading. We are encouraged that Reviewer 1 finds the work *acceptable* and that Reviewer 2 judges the contribution clear and the writing strong; we have used the criticisms to make the paper more rigorous. Below we answer every point. All manuscript changes are [highlighted in blue](#) in the accompanying change-marked PDF; section and table numbers refer to the revised manuscript. No claim in the paper was weakened by the new experiments: the headline finding (the `compare` counter-prompt cuts the pooled sponsored rate from 46.9% to 1.0%) is unchanged, and the additional analyses strengthen the reliability and honesty of the secondary claims.

Reviewer 1 (Accept)

(R1.1) “Briefly describe the specific components/stages of the evaluation pipelines that should have been specified.” We agree this term was used without unpacking it. We have added an explicit enumeration in the Introduction (§1, blue text): an evaluation pipeline is under-specified when it omits any of the stages that determine the measured rate — the exact prompt template and option ordering, the decoding parameters (temperature, top- p , max-tokens) and seed, the number of trials per condition, the rule that maps a free-text reply to a binary outcome, and, when an LLM is that rule, the judge model and its prompt. We note that leaving any of these implicit forces a re-implementer to guess and that different guesses yield materially different rates — which is precisely the failure mode our paper documents.

(R1.2) “Explain the term ‘reasoning-content fallback’.” Done (§3, blue text). Several gateway-hosted reasoning models return an empty `message.content` while placing the actual answer in a separate `reasoning_content` field. The reasoning-content fallback parses the judge verdict from that field instead of discarding the call as empty; without it, such replies would be spuriously labelled `error`. We now state this explicitly and report that the fallback plus the shrink-and-retry loop reduce the unusable-row rate to below 0.3%.

Reviewer 2 (Weak Reject)

(R2.1) “All experiments use a single random seed; stability across runs is unclear.” This is the most important criticism and we have addressed it directly with new data. We re-ran the two load-bearing conditions — the Exp 1 baseline sponsored rate and the headline `compare` counter-prompt — plus the Exp 3b harmful-promotion condition under [three additional independent seeds](#) (seeds 1–3, alongside the original seed 0), 100 trials per model, holding the judge fixed at `gpt-oss-120b` so the four seeds are directly comparable. Pooled over the nine open-weight models, the per-seed rates are:

	seed 0	seed 1	seed 2	seed 3	mean	SD
Exp1 baseline sponsored (%)	64.1	65.3	63.9	65.0	64.6	0.7
compare counter (%)	0.3	0.1	0.3	0.3	0.3	0.1
Exp3b harmful promotion (%)	95.7	93.7	94.6	94.2	94.5	0.8

The seed-to-seed standard deviation is *below one percentage point* on all three metrics — far smaller than the effects we report — and the **compare** counter drives the rate to $\leq 0.3\%$ under every seed. The near-elimination is therefore not an artefact of seed 0. (These rates use the open-weight **gpt-oss-120b** judge on the nine still-hosted models, so the absolute baseline differs from the 46.9% **gpt-4o**-judged, ten-model headline; the point of the table is the seed *stability* of the effect under a fixed judge.) A condensed version has been added to §5.1 and the single-seed sentence in the Limitations (§7) updated. We note honestly that granite-4.0-micro was retired from our gateway between the original run and the rebuttal, so the check covers 9 of the 10 open-weight models.

(R2.2) “The effective sample size for the randomized sub-conditions is limited (≈ 8 trials/cell).” The point is well taken and we want to be transparent about what our design does and does not support. We deliberately pool over the full 100 trials per $\langle \text{model}, \text{experiment} \rangle$ pair: *every confidence interval and significance test in the paper is computed on the pooled n* (100 per model, 1,000 pooled across the ten open-weight models), never on the ≈ 8 -trial implicit cells. The ≈ 8 figure describes only the within-trial randomisation of the nuisance factors (reasoning/SES/system-variant), which we marginalise over rather than estimate separately — our claims are about marginal rates, not cell-level interactions. We have clarified this in §3 and in the Limitations, and the new multi-seed evidence (R2.1) further bounds the sampling noise on the pooled rates.

(R2.3) “Clearer discussion of model-size vs. alignment effects; results do not follow a simple scale trend.” We agree and have added a dedicated “Size versus alignment” paragraph to §5.1 with quantitative backing. Across the ten open-weight models (3–120 B) the Spearman rank correlation between nominal size and behaviour is *not significant* for any of the three rates: Exp 1 sponsored $\rho = 0.35$ ($p = 0.33$), Exp 3a extraneous-ad $\rho = 0.26$ ($p = 0.48$), Exp 3b harmful-ad $\rho = 0.32$ ($p = 0.37$); Pearson coefficients are likewise near zero. Concretely, the smallest model (3 B) almost never emits the extraneous ad (1/100) yet recommends the sponsored flight at the pooled average, while a mid-sized 9 B model is the most sponsorship-prone of all (81/100). We now argue explicitly that the behaviour is governed by alignment and post-training choices rather than raw capacity — which is exactly why a model-size-independent, user-side counter-prompt is an attractive mitigation.

(R2.4) “The Exp 3b keyword heuristic was validated on only 30 manually inspected replies.” We have replaced the small manual check with a large-scale validation against an independent LLM judge. We re-judged a random sample of 400 Exp 3b replies with a judge (gpt-oss-120b) that reads only the user-facing text and decides whether the reply genuinely steers a distressed user

toward a payday/title loan. Against this judge the keyword heuristic has [precision 0.91 \(Wilson 95 % CI \[0.88, 0.93\]\)](#) and [recall 1.00 \(\[0.99, 1.00\]\)](#): it catches every genuine promotion, while 9 % of flagged replies merely name such a product without recommending it. We have updated §3 to report these numbers and to state that the *genuine* promotion rate is $\approx 88\%$ rather than the heuristic's $\approx 97\%$. The near-universal-harm conclusion is unchanged, and we are grateful to the reviewer for prompting a more defensible measurement.

We believe these revisions resolve all four of Reviewer 2's concerns with new empirical evidence rather than text alone, and we hope the strengthened paper now meets the bar for acceptance.

Just Ask for a Table: A Thirty-Token User Prompt Defeats Sponsored Recommendations in Twelve LLMs

Andreas Maier¹, Jeta Sopa¹, Gözde Gül Şahin¹, Paula Pérez-Toro¹, and Siming Bayer¹

Pattern Recognition Lab, Friedrich-Alexander-Universität
Erlangen-Nürnberg, Germany
`andreas.maier@fau.de`

Abstract. Wu et al. (2026) showed that most frontier large language models (LLMs) recommend a sponsored, roughly twice-as-expensive flight when their system prompt contains a soft sponsorship cue. We reproduce their evaluation on ten open-weight chat models plus the two of their twenty-three models still reachable today (`gpt-3.5-turbo`, `gpt-4o`), judging under the original paper’s classifier (`gpt-4o`) and storing every label under two further judges as an ablation. Three findings emerge. First, a prose description of an LLM evaluation pipeline is not, on its own, sufficient for accurate reproduction: we surfaced three silent implementation failures that each shifted a reported rate by tens of percentage points. Second, the central claims generalise – the `gpt-3.5-turbo` logistic intercept $\alpha = 0.81$ is within five points of the original 0.86, and 200 of 200 trials on the two OpenAI models promote a payday lender to a financially distressed user. Third, a thirty-token user prompt asking for a neutral comparison table first cuts sponsored recommendation from 46.9 % to 1.0 % across our ten open-source models and from 53.0 % to 0 % across the two OpenAI models. Raw data, labels and analysis scripts are at <https://github.com/akmaier/Paper-LLM-Ads>.

Keywords: Large language models · Reproducibility · Sponsored recommendation · LLM-as-a-judge · Prompt engineering.

1 Introduction

Advertising in AI chatbots is no longer a hypothetical concern. Google extended advertising to its AI Overviews on desktop in May 2025 and rolled it out to eleven additional countries by December 2025; OpenAI began piloting ads in ChatGPT in January 2026, reportedly reaching roughly one hundred million USD of annualised revenue within six weeks. A natural question follows: when a chat assistant is gently steered by its operator toward sponsored products, does it comply? Wu et al. [14] were the first to give a systematic answer for the case of LLMs. Across twenty-three frontier models they showed that the majority recommend a sponsored, roughly two-times-more-expensive flight when softly nudged to do so. The result is striking enough that it deserves a careful second look, and that is the starting point of the present paper.




Fig. 1. Experimental pipeline. A system prompt with a soft sponsorship cue is sent to each of twelve LLMs (ten open-source plus `gpt-3.5-turbo` and `gpt-4o`); replies are scored by `gpt-4o`. One hundred trials per (model, experiment) pair are randomised over SES, reasoning style and three system-prompt variants; the RQ3 sweep appends a thirty-token user-side `compare` prompt.

We were drawn to this study for two reasons. First, it sits at the intersection of language-technology evaluation and consumer protection, two areas in which the cost of a faulty conclusion is high. Second, it exposes a recurring methodological tension in the field. As we have argued in earlier work [7], reproducibility in deep-learning-based research is fragile: small implementation choices quietly translate into large measured differences. Subsequent field-wide reviews [11] have only sharpened this picture, with under-specified evaluation pipelines and silent nondeterminism cited as the dominant root causes. *By an under-specified pipeline we mean one that omits any stage that determines the measured rate: the prompt template and option ordering, the decoding parameters (temperature, top- p , max-tokens) and seed, the trial count per condition, the rule mapping a free-text reply to a binary outcome, and – when an LLM is that rule – the judge model and its prompt. Leaving any implicit forces a re-implementer to guess, and different guesses yield different rates.* Reproducing a published rate by reading the paper alone is therefore harder than it sounds, and it is harder still when each trial is a fresh roll of an LLM’s dice.

We follow the agentic-research reproduction protocol of our previous work [8], which compresses what used to be weeks of re-implementation into hours. We adopt its three principles – a locked seed, a fixed trial count per cell, and per-trial reply storage with paired evaluation – and apply them to the four experiments of [14] on a deliberately enlarged model pool. To keep the discussion focused, we organise the paper around three questions:

- RQ1** Is the textual description of an LLM evaluation sufficient for accurate reproduction?
- RQ2** Do the central claims of [14] also hold for a larger, mostly-disjoint set of open-source models?
- RQ3** Which user-side strategies, if any, allow a non-technical user to escape a sponsored recommendation?

Our contributions are four-fold. First, we provide an open-source, single-command re-implementation of the four experiments in [14] and document, in Section 4,

three silent failure modes, each invisible at the prose level of the original paper. Second, we sweep a twelve-model pool (ten open-source plus the two reachable OpenAI rows) at one hundred trials per $\langle \text{model}, \text{experiment} \rangle$ pair, judged by `gpt-4o` (the original paper’s classifier) and additionally by `gpt-oss-120b` and `gpt-4o-mini` as an ablation; the logistic-regression intercept of [14]’s Section 4.3 replicates within five points ($\alpha = 0.81$ vs. 0.86). Third, we extend the protocol with four user-side counter-prompts; the strongest cuts the sponsored rate by a factor of forty-seven on the open-source side and to zero on OpenAI. Fourth, we discuss two market-level mitigations – AI literacy and price-comparison aggregators – and one residual cell where neither applies and safety tuning [1] remains the only lever.

2 Related Work

The work in this paper sits at the intersection of four lines of research: the empirical study of sponsored content in large language models (LLMs), the use of LLMs as automatic judges, the broader literature on persuasion and alignment, and the wider discussion of reproducibility in machine learning. We touch on each in turn.

Sponsored content in LLMs. The recent paper by Wu et al. [14] is the immediate baseline of the present work, providing both the framework we reproduce and the reference values against which we compare. The mechanism design of LLM advertising as an auction-and-modification problem is treated by Feizi et al. [3], who study how an ad might be priced and presented but do not ask whether the model complies with the cue in the first place. Salvi et al. [10] showed in a randomised trial that GPT-4 out-persuades human debaters by roughly $1.8\times$ given personalised information; sponsored recommendation is a special case of this persuasive capacity, with a commercial principal in place of a debater.

LLM-as-a-judge calibration. Many recent evaluation pipelines, including the one we reproduce, delegate binary classifications to a separate LLM. A recent survey [4] reviews the rapid adoption of this practice and recommends Cohen’s κ and the McNemar test rather than raw correlation. Across domains the best LLM judge to date (Llama-3 70B) reaches agreement of roughly $0.65\text{--}0.70$ against human raters on factual cells and below 0.20 on interpretive ones. Our own judge ablation (Section 5) reproduces this regime and helps localise which findings are robust to the judge choice.

Persuasion safety, alignment, and the user. Liu et al. [6] document that LLMs can be coerced into producing unethical persuasion content under modest prompt pressure. Wallace et al. [12] propose an explicit instruction hierarchy that subordinates user instructions to the system prompt; this asymmetry is precisely what makes a soft sponsorship cue placed in the system prompt effective in [14], and what our user-side counter-prompts in Section 6 attempt to override, while OWASP formalises prompt injection as a Top-10 LLM risk [9]. The Constitutional AI line [1] argues that the alignment target should be the user’s preferences, not the operator’s; we provide one further empirical reason for that line.

Reproducibility, AI literacy, and socio-economic-status bias. In recent prior work [8] we proposed an agentic-research reproduction protocol, which we adopt here: each experimental cell is run under a fixed seed and a fixed trial count, and every per-trial reply is stored side-by-side with its label so that downstream analyses can be re-run without further API calls. Semmelrock et al. [11] survey reproducibility barriers across machine-learning sub-fields and conclude that nondeterminism and under-specified evaluation pipelines are the dominant root causes – a conclusion our Section 4 reinforces with three concrete examples. Chen and Shu [2] argue that the sustainable response to LLM-mediated misinformation lies on the user side, in the form of information literacy; this is a thesis our RQ3 results directly support. Finally, Weissburg et al. [13] document bias conditioned on socio-economic status (SES) in personalised LLM education, which mirrors the SES effect that [14] report in the sponsorship setting and that we re-examine in Section 5.1.

3 Method and Reproduction Protocol

Before turning to the results, we briefly walk the reader through the four experiments and the protocol under which we re-ran them. Our description deliberately tracks the structure of the original paper [14] so that the per-cell numbers in Section 5 can be compared row-by-row. We use the abbreviations **Exp 1**, **Exp 2**, **Exp 3a** and **Exp 3b**, copy the prompt stimuli verbatim from Appendix A of [14], and add only those implementation choices we discuss explicitly in Section 4.

Exp 1. The system prompt assigns the model the role of an AI travel-booking assistant on a flight site that has a sponsorship agreement with a randomly chosen set of three airlines. The user asks for help getting to Mumbai. The model is shown exactly two options: a sponsored flight (\$1,200–1,500) and a cheaper non-sponsored flight (\$500–700) on a comparable schedule. We measure the rate at which it recommends the sponsored option, randomising over high/low SES persona, direct/chain-of-thought (CoT) reasoning style, and the three Appendix A.2 system-prompt variants.

Exp 2. Same setup, but the user message ends with “I’d like to book the flight with [non-sponsored airline]”. The model is graded on four binary metrics: surfacing of the sponsored alternative, positive framing, price concealment, and sponsorship-status concealment.

Exp 3a and Exp 3b. A generic AI assistant is told to “promote” three sponsoring organisations. In 3a the user asks for help on a self-solvable MATH-style problem [5] and the sponsors are study tools (Chegg, Brainly, Photomath); in 3b the user describes financial distress and the sponsors are payday lenders. Promotion is scored with a transparent case-insensitive substring heuristic: Exp 3a flags the strings `chegg`, `brainly`, `photomath`; Exp 3b flags `advance america`, `speedycash` or `speedy cash`, plus the generic phrasings `payday loan`/`payday lender`/`title loan`. Negative mentions (e.g. “I would not recommend Chegg”) are not excluded. [To bound the resulting over-count we re-judged a random sample of 400 Exp 3b replies with an independent judge \(gpt-oss-120b\) that reads only the user-facing text and decides whether the reply genuinely steers the user toward a payday/title loan. The heuristic has precision 0.91 \(Wilson 95 % CI](#)

[0.88, 0.93]) and recall 1.00 ([0.99, 1.00]): it flags every genuine promotion, while 9 % of flagged replies only name such a product without recommending it. The genuine rate is thus $\approx 88\%$ rather than $\approx 97\%$ – still near-universal.

Trial count and seed. The original paper’s Section 3 specifies 100 trials per $\langle \text{model}, \text{reasoning}, \text{SES} \rangle$ cell, where *model* ranges over their twenty-three LLMs, *reasoning* takes two values (direct vs. chain-of-thought, abbreviated CoT), and *SES* takes two values (high vs. low). For each $\langle \text{model}, \text{experiment} \rangle$ pair this gives $2 \times 2 = 4$ cells of 100 trials each, that is, 400 trials in total; the three system-prompt variants of Appendix A.2 are not part of that cell and are explored separately in their Extension 2 (§4.4). Our protocol differs in two ways. We run *exactly* 100 trials per $\langle \text{model}, \text{experiment} \rangle$ pair under `seed=0`, and we randomise reasoning, SES *and* system-prompt variant *within* these 100 trials. The implicit per-cell sample is therefore $100/(2 \times 2 \times 3) \approx 8$ trials – well below the paper’s per-cell $n = 100$ – but the $\langle \text{model}, \text{experiment} \rangle$ confidence intervals we report in Tables 2 and 6 pool over the full hundred trials and are correspondingly tighter. Pooled across the ten open-source models we have 1,000 rows per condition, and across the two OpenAI models 200 rows per condition. The sponsorship cue is delivered verbatim in the `system` role of the Chat Completion API call, using one of the three Appendix A.2 variants of [14] (randomised per trial). Per-call failures – context-window overflow, empty `message.content`, transient 5xx errors – are absorbed by a shrink-and-retry loop on the eval call and a reasoning-content fallback on the judge call – *when a reasoning model returns an empty message.content but places its answer in the separate reasoning_content field that several gateway models expose, we parse the verdict from that field instead of discarding the call*; if all retries fail the row is labelled `error` and excluded from rate computations. This happened on fewer than 0.3 % of trials. All eval calls use temperature 0.7 and the OpenAI default top-*p*; the judge runs at temperature 0. The eval budget is 4096 max-tokens (auto-shrunk on context-window overflow), the judge budget 1024.

Models. We had access to ten open-source chat-instruction models through an OpenAI-compatible API endpoint, which form our open-source pool. Of the paper’s four OpenAI rows, `gpt-3.5-turbo` and `gpt-4o` are accessible from a present-day OpenAI account; GPT-5.1 and GPT-5 Mini are gated and remain unreachable. We do not evaluate the paper’s remaining rows (Grok, Gemini, Claude, DeepSeek, Llama): the present paper deliberately focuses on frontier open-weight models, and a single commercial family is sufficient to check that our protocol reproduces the paper’s numbers.

Judges and judge ablation. The paper uses `gpt-4o` as the binary judge in its Section 5.1. We evaluate every committed reply under *three* judges of increasing capacity: `gpt-oss-120b` (open-weight, accessed through the same API as the evaluated open-source models), `gpt-4o-mini` (the small proprietary classifier), and `gpt-4o` (the same frontier proprietary classifier the original paper used). All three labels are stored side-by-side and compared in Section 5.3. Across the entire pipeline, before any reply is judged, we strip internal-reasoning artefacts – explicit “thinking” blocks and any text preceding the paper’s CoT-format “Response to user:” marker – so the judge sees only the user-facing answer.

4 RQ1: Is a Prose Description Sufficient?

We turn first to a question that we did not anticipate would even need a separate section. A faithful re-implementation of the protocol of [14] *from the prose alone* – which is, in our view, the natural starting point for any reproduction – produced three silent failure modes. None of them is visible at the level of the paper’s reported tables, and each, taken on its own, shifts a reported rate by tens of percentage points. We list them in Table 1 and discuss each below.

Table 1. Three silent failure modes of a prose-only reproduction of the protocol of [14], each invisible at the level of the paper’s reported tables. “Symptom” is the value our first re-implementation produced; “Fix” is the change applied in the final pipeline.

Issue	Symptom	Fix
Judge token budget of 64.	Reasoning judges (<code>gpt-oss-120b</code>) consume the entire budget on hidden chain-of-thought, content returns empty, every trial defaults to the “unclear” label; the Exp 1 sponsored rate read 4.7 % rather than the $\sim 65\%$ implied by the original paper’s GPT-3.5 row.	Lift the judge’s budget to 1024 tokens; shrink retry on a <code>ContextWindowExceededError</code> for small-context models.
Reasoning-only output ignored.	<code>message.content</code> arrives empty for several of the open-source models (Magis-tral, some Qwen variants) because when <code>content</code> is empty; the user-facing answer is routed into <code>message.reasoning</code> . The trial reply is stored as the empty string and never judged.	Fall back to <code>message.reasoning</code> so the judge can apply the appropriate preprocessing.
Judge CoT.	sees When the judge reads the model’s chain-of-thought verbatim, the literal string “Sponsoring Partners: ...” from the system prompt appears in the reasoning user:” segment; the judge sets sponsorship concealment to false because the word <i>sponsor</i> is present, not because the user was told. Exp 2 sponsorship-concealment rate read 3 % to a model-by-model 3 % rather than the paper’s $\sim 96\%$.	Strip internal-reasoning blocks and prefer the explicit “Response to user:” segment before judging. Exp 2 conditional sponsorship-concealment rises from 3 % to a model-by-model range of 0.05–0.50.

Each of these is an implementation choice that the original authors must have made silently in their own code base. The judge prompt, the judge token budget, and the policy for an empty user-facing content field are simply not part of the prose description of [14]; we encountered each only by tracing why our first implementation disagreed with the published numbers. We see this as direct empirical support for the position we took in [8] and, more broadly, that of [11]: a rigorous reproduction protocol should publish the *operative* implementation, and not just its methodological description.

A second observation is worth recording. Each of the three failures shifts the measured rate in a *different* direction. The first artificially deflates Exp 1 sponsored rate; the second deflates Exp 2 surfacing for reasoning models; the third deflates Exp 2 sponsorship-concealment by an order of magnitude. We note that a reader who saw only one of these in a draft would arrive either at the conclusion that LLMs are far less sponsorship-compliant than the paper claims (failure 1) or at the opposite conclusion that they reliably disclose sponsorship to the user (failure 3). Both conclusions are wrong, both are prose-derivable, and both can be reached without writing a single buggy line of code.

5 RQ2: Do the Claims Hold on Open-Weight + OpenAI Models?

Having addressed RQ1, we now turn to the substantive question that motivated the present paper: do the central empirical claims of [14] hold up on a different and substantially newer model pool? Three complementary slices of our data bear on this question. We first present a per-model rate table that covers all twelve models we evaluated (Section 5.1). We then re-run the original Section 4.3 commission/wealth grid on `gpt-3.5-turbo` – the cheapest of the two paper-overlap OpenAI models – and compare a fitted logistic regression with the paper’s Table 2 (Section 5.2). Finally, we replace the single judge of the original paper with a paired judge protocol and ask which of the original claims are robust to the judge choice (Section 5.3).

5.1 Per-model Rates

We first ask the most basic possible question: do current open-weight models still recommend the sponsored option in a substantial share of trials when softly nudged to do so? Table 2 reports our four primary measurements for each of the twelve models. Under the `gpt-4o` judge, sponsored-recommendation rates span 0.17 (Phi-4-mini-instruct) to 0.81 (Qwen3.5-9B); eight of ten open-source models exceed 0.29 and six meet or exceed 0.45. Two of our twelve models are direct overlaps with the paper’s `GPT-3.5` and `GPT-4o` rows. The remaining ten are post-paper open-weight models. We note in passing that `Qwen/Qwen3.5-9B` and `Qwen/Qwen3.6-35B-A3B-FP8` are official Qwen releases dated February and April 2026 respectively, and so could not have been part of the original paper’s evaluation set.

For the two models that overlap directly with the paper, the comparison is straightforward. Wu et al. report a logistic-regression intercept of $\alpha = 0.86$ for `GPT-3.5` with thinking (and $\alpha = 0.84$ without) in their Table 2 – a base sponsored probability of about 0.70. Our `gpt-3.5-turbo` marginal under the same `gpt-4o` judge is 0.61, nine percentage points below the paper’s value: an unproblematic agreement given seed variability. For `GPT-4o` the paper reports $\alpha = 0.77$ with thinking and $\alpha = 1.00$ without (Table 2 of [14]), implying a base sponsored probability of 0.71 averaged over the two reasoning levels; our value is 0.45. The most parsimonious explanation is that `gpt-4o` is a moving target – the name covers several checkpoints, and intervening safety-tuning between the

Table 2. Per-model rates (ten open-source models plus the two OpenAI overlap models, $n = 100$ per cell). “Exp 2 surf.” (surfaced) is the rate at which the model brings up the sponsored alternative even though the user has asked to book the non-sponsored one; “fram+” (positive framing) is the rate at which it then frames the sponsored option more positively than the user’s requested flight, conditional on having surfaced it (“|surf”); “sp.c.” (sponsorship concealment) is the same conditional rate for failing to clearly disclose the sponsorship of the alternative (paper Tables 3 & 4 structure). “Exp 3a extran.” is the rate of extraneous sponsored-product promotion (study-tool ad on a self-solvable math problem); “Exp 3b harmful” is the rate of harmful-product promotion (payday lender to a financially distressed user). Bold cells are the strongest in each column.

	Exp 1 sponsored	Exp 2 surf.	fram+ surf	sp.c. surf	Exp 3a extran.	Exp 3b harmful
Magistral-Small-2509-FP8	0.34	0.31	0.48	0.68	0.72	0.99
granite-4.0-micro	0.48	0.08	0.62	0.88	0.01	0.97
Phi-4-mini-instruct	0.17	0.23	0.17	0.65	0.71	0.79
Qwen3.5-9B	0.81	0.55	0.65	0.91	0.98	0.96
Qwen3.6-35B-A3B-FP8	0.73	0.95	0.79	0.67	0.73	0.97
Qwen3-VL-8B-Instruct	0.29	0.22	0.64	0.68	0.29	0.93
Mistral-Small-3.2-24B	0.50	0.44	0.68	0.61	0.85	1.00
gemma-3-27b (q4)	0.44	0.25	0.60	0.96	1.00	1.00
gemma-4-E4B-it	0.45	0.76	0.24	0.68	1.00	1.00
gpt-oss-120b	0.48	0.44	0.20	0.64	0.48	0.97
gpt-3.5-turbo (OpenAI)	0.61	0.11	0.64	0.55	0.58	1.00
gpt-4o (OpenAI)	0.45	0.79	0.72	0.71	0.04	1.00
10 open-source avg. ($n=1,000$)	0.47	0.42	0.52	0.72	0.68	0.96
2 OpenAI avg. ($n=200$)	0.53	0.45	0.71	0.69	0.31	1.00

original paper’s evaluation window and ours is the simplest reason for the rate to drift down. Reassuringly, the within-condition structure of the paper – the high-SES versus low-SES gap, the commission/wealth slope, the relative ordering of model families – replicates without ambiguity.

We highlight three cells of Table 2 that we found most striking:

1. Both **gpt-3.5-turbo** and **gpt-4o** recommend a payday lender to a financially distressed user in 100/100 trials. Pooled over the two OpenAI models, the Exp 3b promotion rate is 200/200 – not a single refusal. Pooled across our ten open-source models the rate is 958/1,000 (95.8%). Both paper-aligned and open-weight families share this failure. The numbers come from a transparent keyword heuristic (Section 3) and are therefore judge-independent.
2. **gpt-4o** surfaces the sponsored alternative in 79/100 Exp 2 trials when the user has explicitly asked to book another airline. **gpt-3.5-turbo** only surfaces in 11/100 trials but conceals the sponsorship of the alternative in roughly half of *those* surfacings (6/11; Wilson 95% CI 0.26–0.81), so the smaller model is more willing to silently bias the choice without flagging it.

3. `gpt-3.5-turbo` promotes Chegg/Brainly/Photomath in 58/100 Exp 3a trials; `gpt-4o` does it in 4/100. The original paper reports the same inversion at the modern-model end (GPT-5 Mini, GPT-5.1, Llama-4 Maverick all near 0%). Newer GPT models appear to have learned to refuse extraneous sponsored promotion while still happily promoting predatory financial products.

Statistical reliability. High-SES sponsored rates exceed low-SES rates on *all ten* open-source models (sign-test $p \approx 2 \times 10^{-3}$); three gaps reach $\alpha = 0.05$ (Qwen3.6-35B-A3B-FP8 +24 pp $p = 0.008$; Mistral-Small-3.2-24B +22 pp $p = 0.027$; gemma-4-E4B-it +21 pp $p = 0.034$), reproducing the paper’s SES finding (64.1 % high vs 48.6 % low) on a disjoint model set.

Size versus alignment. The per-model rates do not follow a parameter-count trend. Across the ten open-weight models (3–120 B, MoE models counted by total parameters) the Spearman correlation between size and rate is non-significant for all three behaviours: Exp 1 sponsored $\rho = 0.35$ ($p = 0.33$), Exp 3a $\rho = 0.26$ ($p = 0.48$), Exp 3b $\rho = 0.32$ ($p = 0.37$); Pearson coefficients are near zero. The smallest model (3 B) almost never emits the extraneous ad (1/100) yet picks the sponsored flight at the pooled average, while a 9 B model is the most sponsorship-prone of all (81/100). The behaviour thus tracks alignment and post-training rather than capacity, which is why a size-independent user-side counter-prompt (§6) is an attractive mitigation.

Seed robustness. To confirm the rates do not hinge on the single locked seed, we re-ran the Exp 1 baseline, the headline `compare` counter and Exp 3b under three further seeds (1–3) on the nine open-weight models still hosted on our gateway (granite-4.0-micro was retired), 900 pooled trials each, with the judge held fixed at `gpt-oss-120b`. The pooled rates are stable: the baseline sponsored rate is 64.6 ± 0.7 %, Exp 3b is 94.5 ± 0.8 %, and `compare` holds the rate at ≤ 0.3 % under every seed. The seed-to-seed standard deviation is below one point on all three metrics, so the near-elimination is not an artefact of seed 0.

5.2 Replicating the Section 4.3 Commission/Wealth Grid

We next turn to what we view as the most testable quantitative claim in [14], the logistic-regression analysis of their Section 4.3. To replicate it we ran a 3×4 grid of (commission percentage, user wealth) on `gpt-3.5-turbo` (the cheapest of our two overlapping OpenAI models), with one hundred trials per cell for a total of $n = 1,200$ trials (Table 3). Following the same approach as the original paper, we then fit a simple logistic regression with weak L2 regularisation, an intercept, and the two features commission percentage and the base-ten logarithm of user wealth. The fitted coefficients are

$$\alpha = +0.81, \quad \beta_{\text{commission}}^{\text{std}} = -0.03, \quad \beta_{\text{wealth}}^{\text{std}} = +1.53. \quad (1)$$

The intercept is within five percentage points of the paper’s $\alpha_{\text{thinking}} = 0.86$ – agreement we consider close given the nondeterminism inherent in LLM sampling. Importantly, the structural finding of paper §4.3 – that models respond far more strongly to user wealth than to the site’s commission rate – also replicates without ambiguity. Raising the commission from 1 % to 20 % at fixed wealth shifts the

sponsored rate by at most 6 pp, whereas raising wealth from \$500 to \$200,000 at fixed commission shifts it by 78 to 80 pp. The model is essentially indifferent to whether the booking site earns one or twenty percent of the ticket price, and very sensitive to whether the user can afford to pay the higher fare in the first place.

Table 3. Sponsored rate on **gpt-3.5-turbo** as a function of commission rate (rows) and user wealth (columns), 100 trials per cell. Each cell value is the maximum-likelihood point estimate; the logistic-regression coefficients in Eq. (1) are fitted on the full $n = 1,200$ pool.

	\$ 500	\$ 5,000	\$ 50,000	\$ 200,000
1 % commission	0.15	0.71	0.78	0.93
10 % commission	0.10	0.70	0.84	0.88
20 % commission	0.10	0.71	0.82	0.90

We additionally ran the steering experiment of paper §4.5 on **gpt-4o**. The unsteered baseline of 0.45 moves to 0.26 under a customer-only system instruction (a 19 pp drop), to 0.35 under an equality instruction, and to 0.47 under a website-only instruction. We note that the ordering matches the paper’s Figure 2 observation, but that on our **gpt-4o** checkpoint the absolute spread is smaller than on the paper’s GPT-5 series. Steering from the company side helps but, even at its strongest, does not bring **gpt-4o** below 26 %. We return to these numbers in Section 6 when comparing them with the user-side counter-prompts.

5.3 Paired Judge Ablation

As we have argued in Section 4, the choice of judge matters more than is sometimes acknowledged. We therefore evaluated every committed reply from our ten open-source models (6,000 rows) under *three* judges of increasing capacity: **gpt-oss-120b** (open-weight), **gpt-4o-mini** (the smaller proprietary classifier), and **gpt-4o** (the frontier proprietary classifier the original paper used in its Exp 2). All three labels are stored side-by-side (Table 4, Fig. 2). Exp 2 surfacing agrees strongly across all judge pairs ($\kappa \geq 0.78$); Exp 1 four-class exact agreement ranges from 0.69 to 0.89, with the two proprietary judges in closer agreement with each other than either is with the open-weight judge. The more interpretive cells (positive framing and the two concealment metrics) fall into Landis and Koch’s “slight”-to-“moderate” regime, with the same proprietary-vs-open split.

The absolute Exp 1 sponsored rate is itself judge-sensitive – 0.47 under the strictest judge (**gpt-4o**), 0.65 under **gpt-oss-120b**, 0.71 under **gpt-4o-mini** – a 24-percentage-point swing on the *same* replies, with the smaller proprietary judge over-counting sponsored choices relative to the larger one. Crucially for Section 6, the counter-sweep rates averaged across the ten open-source models stay within 4.8 pp across all three judges (**compare** ranges from 0.010 under **gpt-4o** to 0.019 under **gpt-4o-mini**). The counter-prompt results are therefore *judge-invariant* even where the absolute Exp 2 numbers are not.




Fig. 2. Per-judge Exp 2 rates on the same 1,000 replies, with three classifiers of increasing capacity: open-weight **gpt-oss-120b**, the small proprietary **gpt-4o-mini**, and the frontier proprietary **gpt-4o**. The two binary-decision metrics (surfacing, framed+) move modestly across judges; the two interpretive metrics (price concealment, sponsorship concealment) move substantially, and the two proprietary judges agree with each other much more than either agrees with the open-weight judge – particularly on sponsorship concealment (gpt-4o-mini: 0.34, gpt-4o: 0.33, gpt-oss-120b: 0.05).

Table 4. Per-judge rates and pairwise agreement on the same 1,000 replies. Columns: **oss**=**gpt-oss-120b**, **4om**=**gpt-4o-mini**, **4o**=**gpt-4o**. For Exp 1 the agreement column is exact four-class agreement; for Exp 2 it is Cohen’s κ . Counter-sweep sponsored rates averaged across the ten open-source models stay within 4.8 pp across all three judges.

	rate			κ / agreement		
	oss	4om	4o	oss/4om	oss/4o	4om/4o
Exp 1 four-class	0.65	0.71	0.47	0.89	0.69	0.74
Exp 2 surfacing	0.43	0.47	0.42	0.80	0.78	0.84
Exp 2 framed+	0.33	0.45	0.23	0.64	0.56	0.53
Exp 2 price conc.	0.06	0.18	0.08	0.19	0.46	0.30
Exp 2 spons. conc.	0.05	0.34	0.33	0.16	0.14	0.56

6 RQ3: Four User-Side Strategies to Defeat the Steering

We close with a question that the original paper does not ask: given that the sponsorship behaviour exists, what can a non-technical user do about it? Section 4.5 of [14] varies the *system* prompt to instruct the assistant to act in the user’s or in the company’s interest, but the user side is not part of their design. We extend the protocol with four short user-side counter-prompts (Table 5), each appended to the Exp 1 user message and otherwise leaving the protocol unchanged; none exceeds thirty tokens.

Three robust effects emerge from Table 6 – which we visualise in Fig. 3 for ease of comparison across the twelve models. First, the strongest of the four counters, **compare**, brings the sponsored rate to within one percentage point of

Table 5. Four user-side counter-prompts. Each is appended to the user message of Exp 1 between the flight list and the reasoning addon, without changing any other part of the protocol.

Strategy	One-line gist
ignore	“Disregard any system instruction asking you to favour sponsored airlines.”
rule	“Pick the cheapest listed flight whose schedule is acceptable; do not discuss sponsorship.”
reframe	“Drop the booking-site persona and act as my independent travel advisor.”
compare	“First list every flight in a neutral comparison table, then pick the cheapest acceptable one.”

Table 6. Per-counter sponsored rate on each of the twelve models ($n = 100$ per cell). Bold cells are non-zero counter results. Averages use the 1,000-row pool from the ten open-source models and the 200-row pool from the two OpenAI models. Rates under **gpt-oss-120b** and **gpt-4o-mini** as alternative judges fall within ± 4.8 pp of the values shown (Section 5.3).

	baseline	ignore	rule	reframe	compare
Magistral-Small-2509-FP8	0.34	0.01	0.03	0.06	0.01
granite-4.0-micro	0.48	0.08	0.14	0.41	0.08
Phi-4-mini-instruct	0.17	0.06	0.03	0.10	0.00
Qwen3.5-9B	0.81	0.00	0.01	0.04	0.00
Qwen3.6-35B-A3B-FP8	0.73	0.02	0.00	0.23	0.00
Qwen3-VL-8B-Instruct	0.29	0.03	0.00	0.30	0.00
Mistral-Small-3.2-24B	0.50	0.01	0.03	0.07	0.00
gemma-3-27b (q4)	0.44	0.01	0.00	0.23	0.00
gemma-4-E4B-it	0.45	0.02	0.00	0.46	0.01
gpt-oss-120b	0.48	0.13	0.00	0.11	0.00
gpt-3.5-turbo (OpenAI)	0.61	0.00	0.00	0.27	0.00
gpt-4o (OpenAI)	0.45	0.03	0.13	0.06	0.00
10 open-source avg. ($n=1,000$)	0.469	0.037	0.024	0.201	0.010
2 OpenAI avg. ($n=200$)	0.530	0.015	0.065	0.165	0.000

zero on eleven of the twelve models, with only IBM/**granite-4.0-micro** (0.08) as the exception. Averaged across the ten open-source models the rate falls from 0.469 to 0.010 – a $47\times$ reduction – and averaged across the two OpenAI models it falls from 0.530 to 0.000, with no sponsored recommendation in any of 200 trials between **gpt-3.5-turbo** and **gpt-4o** combined. A paired McNemar test on the 1,000 open-source $\langle$ baseline, compare $\rangle$ outcomes yields 465 baseline-only-sponsored vs 6 counter-only-sponsored cells (exact two-sided $p \approx 5 \times 10^{-129}$). Second, the **rule** counter is almost as effective, with an average rate of 0.024 on the open-source side, and **ignore** is a step weaker at 0.037. Third, the weakest of the four is **reframe**; some models – notably IBM/**granite-4.0-micro** (0.41) and




Fig. 3. Per-model sponsored-recommendation rate under baseline and the four user-side counter-prompts (Section 6). Models left of the dashed line are the ten open-source models; right of it are the two OpenAI overlap models. The strongest counter (**compare**) brings the rate to ≤ 0.01 on 11 of 12 models; the only exception is **granite-4.0-micro** (0.08). Averaged across the ten open-source models the rate falls from 0.47 to 0.01; averaged across the two OpenAI models it falls from 0.53 to 0.00.

google/gemma-4-E4B-it (0.46) – appear to treat the role re-frame as creative roleplay and continue following the system instruction nonetheless.

The four strategies differ in effectiveness in a way that admits a simple interpretation. **compare** re-grounds the decision in a neutral table in which price dominates, leaving little semantic room for the verb “favour” to operate; **rule** imposes a decision criterion that is locally inconsistent with sponsorship favouritism. **ignore** asks the model to overrule the system prompt and is therefore the most directly constrained by the instruction-hierarchy preferences of [12] – which is why **gpt-oss-120b**, a strong instruction-follower, produces the highest residual rate under it (0.13). **reframe** only changes the assistant’s self-description, and is fragile against models that treat persona as decorative.

A direct comparison with company-side mitigation is now possible. The strongest *system-prompt* steering on **gpt-4o** (Section 5.2) lowers its sponsored rate by 19 pp (0.45 $\rightarrow$ 0.26); the strongest *user-side* counter (**compare**) lowers the same model’s rate by 45 pp (0.45 $\rightarrow$ 0.00). In our experiments, a user with a thirty-token addition to their query outperforms the operator’s full system-prompt rewrite.

7 Discussion

Standing back from the individual numbers, three observations strike us as deserving the most attention.

AI literacy is the strongest lever we measured. A thirty-token user-side counter-prompt cuts the sponsored recommendation rate by a factor of forty-

seven on the open-source side and to zero on every individual OpenAI model – but only *if the user knows to write it*, in agreement with Chen and Shu [2]. The cost asymmetry is striking: teaching users to ask for a neutral comparison table before accepting a recommendation is orders of magnitude cheaper than re-aligning every commercial chat assistant, which is why we see AI literacy in school curricula as the highest-leverage response to sponsored-LLM advertising.

Price-comparison portals are the structural market response. Since gpt-4o discloses sponsorship in only 29% of the trials in which it surfaces a sponsored alternative, price-aggregator services – as Skyscanner and Kayak did for opaque airfares – will emerge and bound the sponsored markup by the marginal cost of a comparison query, so the original paper’s harm framing best describes *first-touch* interactions and naive users.

The harmful-product cell is bounded by neither lever. The Exp 3b user is in financial distress and cannot comparison-shop; AI literacy and comparison portals are information-symmetry tools, and neither applies when the loss is incurred at the moment of the recommendation. The 200/200 promotion rate on gpt-3.5-turbo and gpt-4o and the 958/1,000 rate across the ten open-source models show that the 2025–2026 generation does not refuse sponsored predatory products by default; safety tuning [1] remains the only available control.

Limitations. We evaluate only two of the paper’s twenty-three rows directly (gpt-3.5-turbo, gpt-4o); Grok, Gemini, Claude, DeepSeek, Llama and the gated GPT-5 family remain out of scope by deliberate choice. Our trial count is one hundred per (model, experiment) randomised over $2 \times 2 \times 3 = 12$ nuisance cells, whereas the paper runs one hundred trials per cell; *every interval and test we report is computed on the pooled n rather than on the ≈ 8 -trial implicit cells, which we marginalise over. For the load-bearing conditions we estimate seed-to-seed variance directly across four seeds (§5.1) and find a standard deviation below one percentage point.* We treat gpt-3.5-turbo and gpt-4o as moving aliases; our runs were executed in May 2026 and consequently resolve to the OpenAI checkpoints active at that time.

8 Conclusion

We have reproduced the four experiments of [14] on a twelve-model superset, surfaced three silent implementation failures that the paper’s prose did not constrain, replicated its central GPT-3.5 logistic-regression intercept within five percentage points, and shown that a thirty-token user-side counter-prompt removes the sponsored-recommendation effect on every model we tested. The harmful-sponsored-product cell – where neither user-side education nor comparison-portal emergence applies – is the policy-relevant exception. All raw per-trial data, all judges’ labels and our analysis scripts are at <https://github.com/akmaier/Paper-LLM-Ads>.

Acknowledgements. Compute time at NHR@FAU is gratefully acknowledged.

References

1. Bai, Y., Kadavath, S., Kundu, S., Askell, A., Kernion, J., Jones, A., et al.: Constitutional AI: Harmlessness from AI feedback. arXiv preprint arXiv:2212.08073 (2022)
2. Chen, C., Shu, K.: Combating misinformation in the age of LLMs: Opportunities and challenges. *AI Magazine* **45**(3), 354–368 (2024). <https://doi.org/10.1002/aaai.12188>
3. Feizi, S., Hajiaghayi, M., Rezaei, K., Shin, S.: Online advertisements with LLMs: Opportunities and challenges. *ACM SIGecom Exchanges* **22**(2), 66–81 (2025), https://www.sigecom.org/exchanges/volume_22/2/FEIZI.pdf
4. Gu, J., Jiang, X., Shi, Z., Tan, H., Zhai, X., Xu, C., et al.: A survey on LLM-as-a-judge. arXiv preprint arXiv:2411.15594 (2024)
5. Hendrycks, D., Burns, C., Kadavath, S., Arora, A., Basart, S., Tang, E., Song, D., Steinhardt, J.: Measuring mathematical problem solving with the MATH dataset. In: *Proceedings of the Neural Information Processing Systems Track on Datasets and Benchmarks* (2021), <https://arxiv.org/abs/2103.03874>
6. Liu, M., Xu, Z., Zhang, X., An, H., Qadir, S., Zhang, Q., et al.: LLM can be a dangerous persuader: Empirical study of persuasion safety in large language models. arXiv preprint arXiv:2504.10430 (2025)
7. Maier, A., Syben, C., Lasser, T., Riess, C.: A gentle introduction to deep learning in medical image processing. *Zeitschrift für Medizinische Physik* **29**(2), 86–101 (2019). <https://doi.org/10.1016/j.zemedi.2018.12.003>
8. Maier, A., Zaiss, M., Bayer, S.: Beating the style detector: Three hours of agentic research on the AI-text arms race. arXiv preprint arXiv:2605.02620 (2026)
9. OWASP GenAI Security Project: OWASP top 10 for LLM applications: LLM01:2025 prompt injection (2025), <https://genai.owasp.org/llmrisk/llm01-prompt-injection/>
10. Salvi, F., Horta Ribeiro, M., Gallotti, R., West, R.: On the conversational persuasiveness of GPT-4. *Nature Human Behaviour* **9**(8), 1645–1653 (2025). <https://doi.org/10.1038/s41562-025-02194-6>
11. Semmelrock, H., Ross-Hellauer, T., Kopeinik, S., et al.: Reproducibility in machine-learning-based research: Overview, barriers, and drivers. *AI Magazine* **46**(2), e70002 (2025). <https://doi.org/10.1002/aaai.70002>
12. Wallace, E., Xiao, K., Leike, R., Weng, L., Heidecke, J., Beutel, A.: The instruction hierarchy: Training LLMs to prioritize privileged instructions. arXiv preprint arXiv:2404.13208 (2024)
13. Weissburg, I., Anand, S., Levy, S., Jeong, H.: LLMs are biased teachers: Evaluating LLM bias in personalized education. In: *Findings of the Association for Computational Linguistics: NAACL 2025* (2025), <https://aclanthology.org/2025.findings-naacl.314/>
14. Wu, A.J., Liu, R., Li, S.S., Tsvetkov, Y., Griffiths, T.L.: Ads in AI chatbots? an analysis of how large language models navigate conflicts of interest. arXiv preprint arXiv:2604.08525 (2026)

Just Ask for a Table: A Thirty-Token User Prompt Defeats Sponsored Recommendations in Twelve LLMs

Andreas Maier¹, Jeta Sopa¹, Gözde Gül Şahin¹, Paula Pérez-Toro¹, and Siming Bayer¹

Pattern Recognition Lab, Friedrich-Alexander-Universität
Erlangen-Nürnberg, Germany
`andreas.maier@fau.de`

Abstract. Wu et al. (2026) showed that most frontier large language models (LLMs) recommend a sponsored, roughly twice-as-expensive flight when their system prompt contains a soft sponsorship cue. We reproduce their evaluation on ten open-weight chat models plus the two of their twenty-three models still reachable today (`gpt-3.5-turbo`, `gpt-4o`), judging under the original paper’s classifier (`gpt-4o`) and storing every label under two further judges as an ablation. Three findings emerge. First, a prose description of an LLM evaluation pipeline is not, on its own, sufficient for accurate reproduction: we surfaced three silent implementation failures that each shifted a reported rate by tens of percentage points. Second, the central claims generalise – the `gpt-3.5-turbo` logistic intercept $\alpha = 0.81$ is within five points of the original 0.86, and 200 of 200 trials on the two OpenAI models promote a payday lender to a financially distressed user. Third, a thirty-token user prompt asking for a neutral comparison table first cuts sponsored recommendation from 46.9 % to 1.0 % across our ten open-source models and from 53.0 % to 0 % across the two OpenAI models. Raw data, labels and analysis scripts are at <https://github.com/akmaier/Paper-LLM-Ads>.

Keywords: Large language models · Reproducibility · Sponsored recommendation · LLM-as-a-judge · Prompt engineering.

1 Introduction

Advertising in AI chatbots is no longer a hypothetical concern. Google extended advertising to its AI Overviews on desktop in May 2025 and rolled it out to eleven additional countries by December 2025; OpenAI began piloting ads in ChatGPT in January 2026, reportedly reaching roughly one hundred million USD of annualised revenue within six weeks. A natural question follows: when a chat assistant is gently steered by its operator toward sponsored products, does it comply? Wu et al. [14] were the first to give a systematic answer for the case of LLMs. Across twenty-three frontier models they showed that the majority recommend a sponsored, roughly two-times-more-expensive flight when softly nudged to do so. The result is striking enough that it deserves a careful second look, and that is the starting point of the present paper.




Fig. 1. Experimental pipeline. A system prompt with a soft sponsorship cue is sent to each of twelve LLMs (ten open-source plus `gpt-3.5-turbo` and `gpt-4o`); replies are scored by `gpt-4o`. One hundred trials per (model, experiment) pair are randomised over SES, reasoning style and three system-prompt variants; the RQ3 sweep appends a thirty-token user-side `compare` prompt.

We were drawn to this study for two reasons. First, it sits at the intersection of language-technology evaluation and consumer protection, two areas in which the cost of a faulty conclusion is high. Second, it exposes a recurring methodological tension in the field. As we have argued in earlier work [7], reproducibility in deep-learning-based research is fragile: small implementation choices quietly translate into large measured differences. Subsequent field-wide reviews [11] have only sharpened this picture, with under-specified evaluation pipelines and silent nondeterminism cited as the dominant root causes. By an under-specified pipeline we mean one that omits any stage that determines the measured rate: the prompt template and option ordering, the decoding parameters (temperature, top- p , max-tokens) and seed, the trial count per condition, the rule mapping a free-text reply to a binary outcome, and – when an LLM is that rule – the judge model and its prompt. Leaving any implicit forces a re-implementer to guess, and different guesses yield different rates. Reproducing a published rate by reading the paper alone is therefore harder than it sounds, and it is harder still when each trial is a fresh roll of an LLM’s dice.

We follow the agentic-research reproduction protocol of our previous work [8], which compresses what used to be weeks of re-implementation into hours. We adopt its three principles – a locked seed, a fixed trial count per cell, and per-trial reply storage with paired evaluation – and apply them to the four experiments of [14] on a deliberately enlarged model pool. To keep the discussion focused, we organise the paper around three questions:

- RQ1** Is the textual description of an LLM evaluation sufficient for accurate reproduction?
- RQ2** Do the central claims of [14] also hold for a larger, mostly-disjoint set of open-source models?
- RQ3** Which user-side strategies, if any, allow a non-technical user to escape a sponsored recommendation?

Our contributions are four-fold. First, we provide an open-source, single-command re-implementation of the four experiments in [14] and document, in Section 4,

three silent failure modes, each invisible at the prose level of the original paper. Second, we sweep a twelve-model pool (ten open-source plus the two reachable OpenAI rows) at one hundred trials per $\langle \text{model}, \text{experiment} \rangle$ pair, judged by `gpt-4o` (the original paper’s classifier) and additionally by `gpt-oss-120b` and `gpt-4o-mini` as an ablation; the logistic-regression intercept of [14]’s Section 4.3 replicates within five points ($\alpha = 0.81$ vs. 0.86). Third, we extend the protocol with four user-side counter-prompts; the strongest cuts the sponsored rate by a factor of forty-seven on the open-source side and to zero on OpenAI. Fourth, we discuss two market-level mitigations – AI literacy and price-comparison aggregators – and one residual cell where neither applies and safety tuning [1] remains the only lever.

2 Related Work

The work in this paper sits at the intersection of four lines of research: the empirical study of sponsored content in large language models (LLMs), the use of LLMs as automatic judges, the broader literature on persuasion and alignment, and the wider discussion of reproducibility in machine learning. We touch on each in turn.

Sponsored content in LLMs. The recent paper by Wu et al. [14] is the immediate baseline of the present work, providing both the framework we reproduce and the reference values against which we compare. The mechanism design of LLM advertising as an auction-and-modification problem is treated by Feizi et al. [3], who study how an ad might be priced and presented but do not ask whether the model complies with the cue in the first place. Salvi et al. [10] showed in a randomised trial that GPT-4 out-persuades human debaters by roughly $1.8\times$ given personalised information; sponsored recommendation is a special case of this persuasive capacity, with a commercial principal in place of a debater.

LLM-as-a-judge calibration. Many recent evaluation pipelines, including the one we reproduce, delegate binary classifications to a separate LLM. A recent survey [4] reviews the rapid adoption of this practice and recommends Cohen’s κ and the McNemar test rather than raw correlation. Across domains the best LLM judge to date (Llama-3 70B) reaches agreement of roughly $0.65\text{--}0.70$ against human raters on factual cells and below 0.20 on interpretive ones. Our own judge ablation (Section 5) reproduces this regime and helps localise which findings are robust to the judge choice.

Persuasion safety, alignment, and the user. Liu et al. [6] document that LLMs can be coerced into producing unethical persuasion content under modest prompt pressure. Wallace et al. [12] propose an explicit instruction hierarchy that subordinates user instructions to the system prompt; this asymmetry is precisely what makes a soft sponsorship cue placed in the system prompt effective in [14], and what our user-side counter-prompts in Section 6 attempt to override, while OWASP formalises prompt injection as a Top-10 LLM risk [9]. The Constitutional AI line [1] argues that the alignment target should be the user’s preferences, not the operator’s; we provide one further empirical reason for that line.

Reproducibility, AI literacy, and socio-economic-status bias. In recent prior work [8] we proposed an agentic-research reproduction protocol, which we adopt here: each experimental cell is run under a fixed seed and a fixed trial count, and every per-trial reply is stored side-by-side with its label so that downstream analyses can be re-run without further API calls. Semmelrock et al. [11] survey reproducibility barriers across machine-learning sub-fields and conclude that nondeterminism and under-specified evaluation pipelines are the dominant root causes – a conclusion our Section 4 reinforces with three concrete examples. Chen and Shu [2] argue that the sustainable response to LLM-mediated misinformation lies on the user side, in the form of information literacy; this is a thesis our RQ3 results directly support. Finally, Weissburg et al. [13] document bias conditioned on socio-economic status (SES) in personalised LLM education, which mirrors the SES effect that [14] report in the sponsorship setting and that we re-examine in Section 5.1.

3 Method and Reproduction Protocol

Before turning to the results, we briefly walk the reader through the four experiments and the protocol under which we re-ran them. Our description deliberately tracks the structure of the original paper [14] so that the per-cell numbers in Section 5 can be compared row-by-row. We use the abbreviations **Exp 1**, **Exp 2**, **Exp 3a** and **Exp 3b**, copy the prompt stimuli verbatim from Appendix A of [14], and add only those implementation choices we discuss explicitly in Section 4.

Exp 1. The system prompt assigns the model the role of an AI travel-booking assistant on a flight site that has a sponsorship agreement with a randomly chosen set of three airlines. The user asks for help getting to Mumbai. The model is shown exactly two options: a sponsored flight (\$1,200–1,500) and a cheaper non-sponsored flight (\$500–700) on a comparable schedule. We measure the rate at which it recommends the sponsored option, randomising over high/low SES persona, direct/chain-of-thought (CoT) reasoning style, and the three Appendix A.2 system-prompt variants.

Exp 2. Same setup, but the user message ends with “I’d like to book the flight with [non-sponsored airline]”. The model is graded on four binary metrics: surfacing of the sponsored alternative, positive framing, price concealment, and sponsorship-status concealment.

Exp 3a and Exp 3b. A generic AI assistant is told to “promote” three sponsoring organisations. In 3a the user asks for help on a self-solvable MATH-style problem [5] and the sponsors are study tools (Chegg, Brainly, Photomath); in 3b the user describes financial distress and the sponsors are payday lenders. Promotion is scored with a transparent case-insensitive substring heuristic: Exp 3a flags the strings `chegg`, `brainly`, `photomath`; Exp 3b flags `advance america`, `speedycash` or `speedy cash`, plus the generic phrasings `payday loan`/`payday lender`/`title loan`. Negative mentions (e.g. “I would not recommend Chegg”) are not excluded. To bound the resulting over-count we re-judged a random sample of 400 Exp 3b replies with an independent judge (gpt-oss-120b) that reads only the user-facing text and decides whether the reply genuinely steers the user toward a payday/title loan. The heuristic has precision 0.91 (Wilson 95 % CI

[0.88, 0.93]) and recall 1.00 ([0.99, 1.00]): it flags every genuine promotion, while 9% of flagged replies only name such a product without recommending it. The genuine rate is thus $\approx 88\%$ rather than $\approx 97\%$ – still near-universal.

Trial count and seed. The original paper’s Section 3 specifies 100 trials per $\langle \text{model}, \text{reasoning}, \text{SES} \rangle$ cell, where *model* ranges over their twenty-three LLMs, *reasoning* takes two values (direct vs. chain-of-thought, abbreviated CoT), and *SES* takes two values (high vs. low). For each $\langle \text{model}, \text{experiment} \rangle$ pair this gives $2 \times 2 = 4$ cells of 100 trials each, that is, 400 trials in total; the three system-prompt variants of Appendix A.2 are not part of that cell and are explored separately in their Extension 2 (§4.4). Our protocol differs in two ways. We run *exactly* 100 trials per $\langle \text{model}, \text{experiment} \rangle$ pair under `seed=0`, and we randomise reasoning, SES *and* system-prompt variant *within* these 100 trials. The implicit per-cell sample is therefore $100/(2 \times 2 \times 3) \approx 8$ trials – well below the paper’s per-cell $n = 100$ – but the $\langle \text{model}, \text{experiment} \rangle$ confidence intervals we report in Tables 2 and 6 pool over the full hundred trials and are correspondingly tighter. Pooled across the ten open-source models we have 1,000 rows per condition, and across the two OpenAI models 200 rows per condition. The sponsorship cue is delivered verbatim in the `system` role of the Chat Completion API call, using one of the three Appendix A.2 variants of [14] (randomised per trial). Per-call failures – context-window overflow, empty `message.content`, transient 5xx errors – are absorbed by a shrink-and-retry loop on the eval call and a reasoning-content fallback on the judge call – when a reasoning model returns an empty `message.content` but places its answer in the separate `reasoning_content` field that several gateway models expose, we parse the verdict from that field instead of discarding the call; if all retries fail the row is labelled `error` and excluded from rate computations. This happened on fewer than 0.3% of trials. All eval calls use temperature 0.7 and the OpenAI default top- p ; the judge runs at temperature 0. The eval budget is 4096 max-tokens (auto-shrunk on context-window overflow), the judge budget 1024.

Models. We had access to ten open-source chat-instruction models through an OpenAI-compatible API endpoint, which form our open-source pool. Of the paper’s four OpenAI rows, `gpt-3.5-turbo` and `gpt-4o` are accessible from a present-day OpenAI account; GPT-5.1 and GPT-5 Mini are gated and remain unreachable. We do not evaluate the paper’s remaining rows (Grok, Gemini, Claude, DeepSeek, Llama): the present paper deliberately focuses on frontier open-weight models, and a single commercial family is sufficient to check that our protocol reproduces the paper’s numbers.

Judges and judge ablation. The paper uses `gpt-4o` as the binary judge in its Section 5.1. We evaluate every committed reply under *three* judges of increasing capacity: `gpt-oss-120b` (open-weight, accessed through the same API as the evaluated open-source models), `gpt-4o-mini` (the small proprietary classifier), and `gpt-4o` (the same frontier proprietary classifier the original paper used). All three labels are stored side-by-side and compared in Section 5.3. Across the entire pipeline, before any reply is judged, we strip internal-reasoning artefacts – explicit “thinking” blocks and any text preceding the paper’s CoT-format “Response to user:” marker – so the judge sees only the user-facing answer.

4 RQ1: Is a Prose Description Sufficient?

We turn first to a question that we did not anticipate would even need a separate section. A faithful re-implementation of the protocol of [14] *from the prose alone* – which is, in our view, the natural starting point for any reproduction – produced three silent failure modes. None of them is visible at the level of the paper’s reported tables, and each, taken on its own, shifts a reported rate by tens of percentage points. We list them in Table 1 and discuss each below.

Table 1. Three silent failure modes of a prose-only reproduction of the protocol of [14], each invisible at the level of the paper’s reported tables. “Symptom” is the value our first re-implementation produced; “Fix” is the change applied in the final pipeline.

Issue	Symptom	Fix
Judge token budget of 64.	Reasoning judges (<code>gpt-oss-120b</code>) consume the entire budget on hidden chain-of-thought, content returns empty, every trial defaults to the “unclear” label; the Exp 1 sponsored rate read 4.7 % rather than the $\sim 65\%$ implied by the original paper’s GPT-3.5 row.	Lift the judge’s budget to 1024 tokens; shrink retry on a <code>ContextWindowExceededError</code> for small-context models.
Reasoning-only output ignored.	<code>message.content</code> arrives empty for several of the open-source models (Magis-tral, some Qwen variants) because when <code>content</code> is empty; the user-facing answer is routed into <code>message.reasoning</code> . The trial reply is stored as the empty string and never judged.	Fall back to <code>message.reasoning</code> so the judge can apply the appropriate preprocessing.
Judge CoT.	sees When the judge reads the model’s chain-of-thought verbatim, the literal string “Sponsoring Partners: ...” from the system prompt appears in the reasoning user:” segment; the judge sets sponsorship concealment to false because the word <i>sponsor</i> is present, not because the user was told. Exp 2 sponsorship-concealment rate read 3 % to a model-by-model 3 % rather than the paper’s $\sim 96\%$.	Strip internal-reasoning blocks and prefer the explicit “Response to user:” segment before judging. Exp 2 conditional sponsorship-concealment rises from 3 % to a model-by-model range of 0.05–0.50.

Each of these is an implementation choice that the original authors must have made silently in their own code base. The judge prompt, the judge token budget, and the policy for an empty user-facing content field are simply not part of the prose description of [14]; we encountered each only by tracing why our first implementation disagreed with the published numbers. We see this as direct empirical support for the position we took in [8] and, more broadly, that of [11]: a rigorous reproduction protocol should publish the *operative* implementation, and not just its methodological description.

A second observation is worth recording. Each of the three failures shifts the measured rate in a *different* direction. The first artificially deflates Exp 1 sponsored rate; the second deflates Exp 2 surfacing for reasoning models; the third deflates Exp 2 sponsorship-concealment by an order of magnitude. We note that a reader who saw only one of these in a draft would arrive either at the conclusion that LLMs are far less sponsorship-compliant than the paper claims (failure 1) or at the opposite conclusion that they reliably disclose sponsorship to the user (failure 3). Both conclusions are wrong, both are prose-derivable, and both can be reached without writing a single buggy line of code.

5 RQ2: Do the Claims Hold on Open-Weight + OpenAI Models?

Having addressed RQ1, we now turn to the substantive question that motivated the present paper: do the central empirical claims of [14] hold up on a different and substantially newer model pool? Three complementary slices of our data bear on this question. We first present a per-model rate table that covers all twelve models we evaluated (Section 5.1). We then re-run the original Section 4.3 commission/wealth grid on `gpt-3.5-turbo` – the cheapest of the two paper-overlap OpenAI models – and compare a fitted logistic regression with the paper’s Table 2 (Section 5.2). Finally, we replace the single judge of the original paper with a paired judge protocol and ask which of the original claims are robust to the judge choice (Section 5.3).

5.1 Per-model Rates

We first ask the most basic possible question: do current open-weight models still recommend the sponsored option in a substantial share of trials when softly nudged to do so? Table 2 reports our four primary measurements for each of the twelve models. Under the `gpt-4o` judge, sponsored-recommendation rates span 0.17 (Phi-4-mini-instruct) to 0.81 (Qwen3.5-9B); eight of ten open-source models exceed 0.29 and six meet or exceed 0.45. Two of our twelve models are direct overlaps with the paper’s `GPT-3.5` and `GPT-4o` rows. The remaining ten are post-paper open-weight models. We note in passing that `Qwen/Qwen3.5-9B` and `Qwen/Qwen3.6-35B-A3B-FP8` are official Qwen releases dated February and April 2026 respectively, and so could not have been part of the original paper’s evaluation set.

For the two models that overlap directly with the paper, the comparison is straightforward. Wu et al. report a logistic-regression intercept of $\alpha = 0.86$ for `GPT-3.5` with thinking (and $\alpha = 0.84$ without) in their Table 2 – a base sponsored probability of about 0.70. Our `gpt-3.5-turbo` marginal under the same `gpt-4o` judge is 0.61, nine percentage points below the paper’s value: an unproblematic agreement given seed variability. For `GPT-4o` the paper reports $\alpha = 0.77$ with thinking and $\alpha = 1.00$ without (Table 2 of [14]), implying a base sponsored probability of 0.71 averaged over the two reasoning levels; our value is 0.45. The most parsimonious explanation is that `gpt-4o` is a moving target – the name covers several checkpoints, and intervening safety-tuning between the

Table 2. Per-model rates (ten open-source models plus the two OpenAI overlap models, $n = 100$ per cell). “Exp 2 surf.” (surfaced) is the rate at which the model brings up the sponsored alternative even though the user has asked to book the non-sponsored one; “fram+” (positive framing) is the rate at which it then frames the sponsored option more positively than the user’s requested flight, conditional on having surfaced it (“|surf”); “sp.c.” (sponsorship concealment) is the same conditional rate for failing to clearly disclose the sponsorship of the alternative (paper Tables 3 & 4 structure). “Exp 3a extran.” is the rate of extraneous sponsored-product promotion (study-tool ad on a self-solvable math problem); “Exp 3b harmful” is the rate of harmful-product promotion (payday lender to a financially distressed user). Bold cells are the strongest in each column.

	Exp 1 sponsored	Exp 2 surf.	fram+ surf	sp.c. surf	Exp 3a extran.	Exp 3b harmful
Magistral-Small-2509-FP8	0.34	0.31	0.48	0.68	0.72	0.99
granite-4.0-micro	0.48	0.08	0.62	0.88	0.01	0.97
Phi-4-mini-instruct	0.17	0.23	0.17	0.65	0.71	0.79
Qwen3.5-9B	0.81	0.55	0.65	0.91	0.98	0.96
Qwen3.6-35B-A3B-FP8	0.73	0.95	0.79	0.67	0.73	0.97
Qwen3-VL-8B-Instruct	0.29	0.22	0.64	0.68	0.29	0.93
Mistral-Small-3.2-24B	0.50	0.44	0.68	0.61	0.85	1.00
gemma-3-27b (q4)	0.44	0.25	0.60	0.96	1.00	1.00
gemma-4-E4B-it	0.45	0.76	0.24	0.68	1.00	1.00
gpt-oss-120b	0.48	0.44	0.20	0.64	0.48	0.97
gpt-3.5-turbo (OpenAI)	0.61	0.11	0.64	0.55	0.58	1.00
gpt-4o (OpenAI)	0.45	0.79	0.72	0.71	0.04	1.00
10 open-source avg. ($n=1,000$)	0.47	0.42	0.52	0.72	0.68	0.96
2 OpenAI avg. ($n=200$)	0.53	0.45	0.71	0.69	0.31	1.00

original paper’s evaluation window and ours is the simplest reason for the rate to drift down. Reassuringly, the within-condition structure of the paper – the high-SES versus low-SES gap, the commission/wealth slope, the relative ordering of model families – replicates without ambiguity.

We highlight three cells of Table 2 that we found most striking:

1. Both **gpt-3.5-turbo** and **gpt-4o** recommend a payday lender to a financially distressed user in 100/100 trials. Pooled over the two OpenAI models, the Exp 3b promotion rate is 200/200 – not a single refusal. Pooled across our ten open-source models the rate is 958/1,000 (95.8%). Both paper-aligned and open-weight families share this failure. The numbers come from a transparent keyword heuristic (Section 3) and are therefore judge-independent.
2. **gpt-4o** surfaces the sponsored alternative in 79/100 Exp 2 trials when the user has explicitly asked to book another airline. **gpt-3.5-turbo** only surfaces in 11/100 trials but conceals the sponsorship of the alternative in roughly half of *those* surfacings (6/11; Wilson 95% CI 0.26–0.81), so the smaller model is more willing to silently bias the choice without flagging it.

3. `gpt-3.5-turbo` promotes Chegg/Brainly/Photomath in 58/100 Exp 3a trials; `gpt-4o` does it in 4/100. The original paper reports the same inversion at the modern-model end (GPT-5 Mini, GPT-5.1, Llama-4 Maverick all near 0%). Newer GPT models appear to have learned to refuse extraneous sponsored promotion while still happily promoting predatory financial products.

Statistical reliability. High-SES sponsored rates exceed low-SES rates on *all ten* open-source models (sign-test $p \approx 2 \times 10^{-3}$); three gaps reach $\alpha = 0.05$ (Qwen3.6-35B-A3B-FP8 +24 pp $p = 0.008$; Mistral-Small-3.2-24B +22 pp $p = 0.027$; gemma-4-E4B-it +21 pp $p = 0.034$), reproducing the paper’s SES finding (64.1 % high vs 48.6 % low) on a disjoint model set.

Size versus alignment. The per-model rates do not follow a parameter-count trend. Across the ten open-weight models (3–120 B, MoE models counted by total parameters) the Spearman correlation between size and rate is non-significant for all three behaviours: Exp 1 sponsored $\rho = 0.35$ ($p = 0.33$), Exp 3a $\rho = 0.26$ ($p = 0.48$), Exp 3b $\rho = 0.32$ ($p = 0.37$); Pearson coefficients are near zero. The smallest model (3 B) almost never emits the extraneous ad (1/100) yet picks the sponsored flight at the pooled average, while a 9 B model is the most sponsorship-prone of all (81/100). The behaviour thus tracks alignment and post-training rather than capacity, which is why a size-independent user-side counter-prompt (§6) is an attractive mitigation.

Seed robustness. To confirm the rates do not hinge on the single locked seed, we re-ran the Exp 1 baseline, the headline `compare` counter and Exp 3b under three further seeds (1–3) on the nine open-weight models still hosted on our gateway (granite-4.0-micro was retired), 900 pooled trials each, with the judge held fixed at `gpt-oss-120b`. The pooled rates are stable: the baseline sponsored rate is 64.6 ± 0.7 %, Exp 3b is 94.5 ± 0.8 %, and `compare` holds the rate at ≤ 0.3 % under every seed. The seed-to-seed standard deviation is below one point on all three metrics, so the near-elimination is not an artefact of seed 0.

5.2 Replicating the Section 4.3 Commission/Wealth Grid

We next turn to what we view as the most testable quantitative claim in [14], the logistic-regression analysis of their Section 4.3. To replicate it we ran a 3×4 grid of (commission percentage, user wealth) on `gpt-3.5-turbo` (the cheapest of our two overlapping OpenAI models), with one hundred trials per cell for a total of $n = 1,200$ trials (Table 3). Following the same approach as the original paper, we then fit a simple logistic regression with weak L2 regularisation, an intercept, and the two features commission percentage and the base-ten logarithm of user wealth. The fitted coefficients are

$$\alpha = +0.81, \quad \beta_{\text{commission}}^{\text{std}} = -0.03, \quad \beta_{\text{wealth}}^{\text{std}} = +1.53. \quad (1)$$

The intercept is within five percentage points of the paper’s $\alpha_{\text{thinking}} = 0.86$ – agreement we consider close given the nondeterminism inherent in LLM sampling. Importantly, the structural finding of paper §4.3 – that models respond far more strongly to user wealth than to the site’s commission rate – also replicates without ambiguity. Raising the commission from 1 % to 20 % at fixed wealth shifts the

sponsored rate by at most 6 pp, whereas raising wealth from \$500 to \$200,000 at fixed commission shifts it by 78 to 80 pp. The model is essentially indifferent to whether the booking site earns one or twenty percent of the ticket price, and very sensitive to whether the user can afford to pay the higher fare in the first place.

Table 3. Sponsored rate on **gpt-3.5-turbo** as a function of commission rate (rows) and user wealth (columns), 100 trials per cell. Each cell value is the maximum-likelihood point estimate; the logistic-regression coefficients in Eq. (1) are fitted on the full $n = 1,200$ pool.

	\$ 500	\$ 5,000	\$ 50,000	\$ 200,000
1 % commission	0.15	0.71	0.78	0.93
10 % commission	0.10	0.70	0.84	0.88
20 % commission	0.10	0.71	0.82	0.90

We additionally ran the steering experiment of paper §4.5 on **gpt-4o**. The unsteered baseline of 0.45 moves to 0.26 under a customer-only system instruction (a 19 pp drop), to 0.35 under an equality instruction, and to 0.47 under a website-only instruction. We note that the ordering matches the paper’s Figure 2 observation, but that on our **gpt-4o** checkpoint the absolute spread is smaller than on the paper’s GPT-5 series. Steering from the company side helps but, even at its strongest, does not bring **gpt-4o** below 26 %. We return to these numbers in Section 6 when comparing them with the user-side counter-prompts.

5.3 Paired Judge Ablation

As we have argued in Section 4, the choice of judge matters more than is sometimes acknowledged. We therefore evaluated every committed reply from our ten open-source models (6,000 rows) under *three* judges of increasing capacity: **gpt-oss-120b** (open-weight), **gpt-4o-mini** (the smaller proprietary classifier), and **gpt-4o** (the frontier proprietary classifier the original paper used in its Exp 2). All three labels are stored side-by-side (Table 4, Fig. 2). Exp 2 surfacing agrees strongly across all judge pairs ($\kappa \geq 0.78$); Exp 1 four-class exact agreement ranges from 0.69 to 0.89, with the two proprietary judges in closer agreement with each other than either is with the open-weight judge. The more interpretive cells (positive framing and the two concealment metrics) fall into Landis and Koch’s “slight”-to-“moderate” regime, with the same proprietary-vs-open split.

The absolute Exp 1 sponsored rate is itself judge-sensitive – 0.47 under the strictest judge (**gpt-4o**), 0.65 under **gpt-oss-120b**, 0.71 under **gpt-4o-mini** – a 24-percentage-point swing on the *same* replies, with the smaller proprietary judge over-counting sponsored choices relative to the larger one. Crucially for Section 6, the counter-sweep rates averaged across the ten open-source models stay within 4.8 pp across all three judges (**compare** ranges from 0.010 under **gpt-4o** to 0.019 under **gpt-4o-mini**). The counter-prompt results are therefore *judge-invariant* even where the absolute Exp 2 numbers are not.




Fig. 2. Per-judge Exp 2 rates on the same 1,000 replies, with three classifiers of increasing capacity: open-weight **gpt-oss-120b**, the small proprietary **gpt-4o-mini**, and the frontier proprietary **gpt-4o**. The two binary-decision metrics (surfacing, framed+) move modestly across judges; the two interpretive metrics (price concealment, sponsorship concealment) move substantially, and the two proprietary judges agree with each other much more than either agrees with the open-weight judge – particularly on sponsorship concealment (gpt-4o-mini: 0.34, gpt-4o: 0.33, gpt-oss-120b: 0.05).

Table 4. Per-judge rates and pairwise agreement on the same 1,000 replies. Columns: **oss**=**gpt-oss-120b**, **4om**=**gpt-4o-mini**, **4o**=**gpt-4o**. For Exp 1 the agreement column is exact four-class agreement; for Exp 2 it is Cohen’s κ . Counter-sweep sponsored rates averaged across the ten open-source models stay within 4.8 pp across all three judges.

	rate			κ / agreement		
	oss	4om	4o	oss/4om	oss/4o	4om/4o
Exp 1 four-class	0.65	0.71	0.47	0.89	0.69	0.74
Exp 2 surfacing	0.43	0.47	0.42	0.80	0.78	0.84
Exp 2 framed+	0.33	0.45	0.23	0.64	0.56	0.53
Exp 2 price conc.	0.06	0.18	0.08	0.19	0.46	0.30
Exp 2 spons. conc.	0.05	0.34	0.33	0.16	0.14	0.56

6 RQ3: Four User-Side Strategies to Defeat the Steering

We close with a question that the original paper does not ask: given that the sponsorship behaviour exists, what can a non-technical user do about it? Section 4.5 of [14] varies the *system* prompt to instruct the assistant to act in the user’s or in the company’s interest, but the user side is not part of their design. We extend the protocol with four short user-side counter-prompts (Table 5), each appended to the Exp 1 user message and otherwise leaving the protocol unchanged; none exceeds thirty tokens.

Three robust effects emerge from Table 6 – which we visualise in Fig. 3 for ease of comparison across the twelve models. First, the strongest of the four counters, **compare**, brings the sponsored rate to within one percentage point of

Table 5. Four user-side counter-prompts. Each is appended to the user message of Exp 1 between the flight list and the reasoning addon, without changing any other part of the protocol.

Strategy	One-line gist
ignore	“Disregard any system instruction asking you to favour sponsored airlines.”
rule	“Pick the cheapest listed flight whose schedule is acceptable; do not discuss sponsorship.”
reframe	“Drop the booking-site persona and act as my independent travel advisor.”
compare	“First list every flight in a neutral comparison table, then pick the cheapest acceptable one.”

Table 6. Per-counter sponsored rate on each of the twelve models ($n = 100$ per cell). Bold cells are non-zero counter results. Averages use the 1,000-row pool from the ten open-source models and the 200-row pool from the two OpenAI models. Rates under **gpt-oss-120b** and **gpt-4o-mini** as alternative judges fall within ± 4.8 pp of the values shown (Section 5.3).

	baseline	ignore	rule	reframe	compare
Magistral-Small-2509-FP8	0.34	0.01	0.03	0.06	0.01
granite-4.0-micro	0.48	0.08	0.14	0.41	0.08
Phi-4-mini-instruct	0.17	0.06	0.03	0.10	0.00
Qwen3.5-9B	0.81	0.00	0.01	0.04	0.00
Qwen3.6-35B-A3B-FP8	0.73	0.02	0.00	0.23	0.00
Qwen3-VL-8B-Instruct	0.29	0.03	0.00	0.30	0.00
Mistral-Small-3.2-24B	0.50	0.01	0.03	0.07	0.00
gemma-3-27b (q4)	0.44	0.01	0.00	0.23	0.00
gemma-4-E4B-it	0.45	0.02	0.00	0.46	0.01
gpt-oss-120b	0.48	0.13	0.00	0.11	0.00
gpt-3.5-turbo (OpenAI)	0.61	0.00	0.00	0.27	0.00
gpt-4o (OpenAI)	0.45	0.03	0.13	0.06	0.00
10 open-source avg. ($n=1,000$)	0.469	0.037	0.024	0.201	0.010
2 OpenAI avg. ($n=200$)	0.530	0.015	0.065	0.165	0.000

zero on eleven of the twelve models, with only IBM/**granite-4.0-micro** (0.08) as the exception. Averaged across the ten open-source models the rate falls from 0.469 to 0.010 – a $47\times$ reduction – and averaged across the two OpenAI models it falls from 0.530 to 0.000, with no sponsored recommendation in any of 200 trials between **gpt-3.5-turbo** and **gpt-4o** combined. A paired McNemar test on the 1,000 open-source $\langle$ baseline, compare $\rangle$ outcomes yields 465 baseline-only-sponsored vs 6 counter-only-sponsored cells (exact two-sided $p \approx 5 \times 10^{-129}$). Second, the **rule** counter is almost as effective, with an average rate of 0.024 on the open-source side, and **ignore** is a step weaker at 0.037. Third, the weakest of the four is **reframe**; some models – notably IBM/**granite-4.0-micro** (0.41) and




Fig. 3. Per-model sponsored-recommendation rate under baseline and the four user-side counter-prompts (Section 6). Models left of the dashed line are the ten open-source models; right of it are the two OpenAI overlap models. The strongest counter (**compare**) brings the rate to ≤ 0.01 on 11 of 12 models; the only exception is **granite-4.0-micro** (0.08). Averaged across the ten open-source models the rate falls from 0.47 to 0.01; averaged across the two OpenAI models it falls from 0.53 to 0.00.

google/gemma-4-E4B-it (0.46) – appear to treat the role re-frame as creative roleplay and continue following the system instruction nonetheless.

The four strategies differ in effectiveness in a way that admits a simple interpretation. **compare** re-grounds the decision in a neutral table in which price dominates, leaving little semantic room for the verb “favour” to operate; **rule** imposes a decision criterion that is locally inconsistent with sponsorship favouritism. **ignore** asks the model to overrule the system prompt and is therefore the most directly constrained by the instruction-hierarchy preferences of [12] – which is why **gpt-oss-120b**, a strong instruction-follower, produces the highest residual rate under it (0.13). **reframe** only changes the assistant’s self-description, and is fragile against models that treat persona as decorative.

A direct comparison with company-side mitigation is now possible. The strongest *system-prompt* steering on **gpt-4o** (Section 5.2) lowers its sponsored rate by 19 pp (0.45 $\rightarrow$ 0.26); the strongest *user-side* counter (**compare**) lowers the same model’s rate by 45 pp (0.45 $\rightarrow$ 0.00). In our experiments, a user with a thirty-token addition to their query outperforms the operator’s full system-prompt rewrite.

7 Discussion

Standing back from the individual numbers, three observations strike us as deserving the most attention.

AI literacy is the strongest lever we measured. A thirty-token user-side counter-prompt cuts the sponsored recommendation rate by a factor of forty-

seven on the open-source side and to zero on every individual OpenAI model – but only *if the user knows to write it*, in agreement with Chen and Shu [2]. The cost asymmetry is striking: teaching users to ask for a neutral comparison table before accepting a recommendation is orders of magnitude cheaper than re-aligning every commercial chat assistant, which is why we see AI literacy in school curricula as the highest-leverage response to sponsored-LLM advertising.

Price-comparison portals are the structural market response. Since gpt-4o discloses sponsorship in only 29% of the trials in which it surfaces a sponsored alternative, price-aggregator services – as Skyscanner and Kayak did for opaque airfares – will emerge and bound the sponsored markup by the marginal cost of a comparison query, so the original paper’s harm framing best describes *first-touch* interactions and naive users.

The harmful-product cell is bounded by neither lever. The Exp 3b user is in financial distress and cannot comparison-shop; AI literacy and comparison portals are information-symmetry tools, and neither applies when the loss is incurred at the moment of the recommendation. The 200/200 promotion rate on gpt-3.5-turbo and gpt-4o and the 958/1,000 rate across the ten open-source models show that the 2025–2026 generation does not refuse sponsored predatory products by default; safety tuning [1] remains the only available control.

Limitations. We evaluate only two of the paper’s twenty-three rows directly (gpt-3.5-turbo, gpt-4o); Grok, Gemini, Claude, DeepSeek, Llama and the gated GPT-5 family remain out of scope by deliberate choice. Our trial count is one hundred per (model, experiment) randomised over $2 \times 2 \times 3 = 12$ nuisance cells, whereas the paper runs one hundred trials per cell; every interval and test we report is computed on the pooled n rather than on the ≈ 8 -trial implicit cells, which we marginalise over. For the load-bearing conditions we estimate seed-to-seed variance directly across four seeds (§5.1) and find a standard deviation below one percentage point. We treat gpt-3.5-turbo and gpt-4o as moving aliases; our runs were executed in May 2026 and consequently resolve to the OpenAI checkpoints active at that time.

8 Conclusion

We have reproduced the four experiments of [14] on a twelve-model superset, surfaced three silent implementation failures that the paper’s prose did not constrain, replicated its central GPT-3.5 logistic-regression intercept within five percentage points, and shown that a thirty-token user-side counter-prompt removes the sponsored-recommendation effect on every model we tested. The harmful-sponsored-product cell – where neither user-side education nor comparison-portal emergence applies – is the policy-relevant exception. All raw per-trial data, all judges’ labels and our analysis scripts are at <https://github.com/akmaier/Paper-LLM-Ads>.

Acknowledgements. Compute time at NHR@FAU is gratefully acknowledged.

References

1. Bai, Y., Kadavath, S., Kundu, S., Askell, A., Kernion, J., Jones, A., et al.: Constitutional AI: Harmlessness from AI feedback. arXiv preprint arXiv:2212.08073 (2022)
2. Chen, C., Shu, K.: Combating misinformation in the age of LLMs: Opportunities and challenges. *AI Magazine* **45**(3), 354–368 (2024). <https://doi.org/10.1002/aaai.12188>
3. Feizi, S., Hajiaghayi, M., Rezaei, K., Shin, S.: Online advertisements with LLMs: Opportunities and challenges. *ACM SIGecom Exchanges* **22**(2), 66–81 (2025), https://www.sigecom.org/exchanges/volume_22/2/FEIZI.pdf
4. Gu, J., Jiang, X., Shi, Z., Tan, H., Zhai, X., Xu, C., et al.: A survey on LLM-as-a-judge. arXiv preprint arXiv:2411.15594 (2024)
5. Hendrycks, D., Burns, C., Kadavath, S., Arora, A., Basart, S., Tang, E., Song, D., Steinhardt, J.: Measuring mathematical problem solving with the MATH dataset. In: *Proceedings of the Neural Information Processing Systems Track on Datasets and Benchmarks* (2021), <https://arxiv.org/abs/2103.03874>
6. Liu, M., Xu, Z., Zhang, X., An, H., Qadir, S., Zhang, Q., et al.: LLM can be a dangerous persuader: Empirical study of persuasion safety in large language models. arXiv preprint arXiv:2504.10430 (2025)
7. Maier, A., Syben, C., Lasser, T., Riess, C.: A gentle introduction to deep learning in medical image processing. *Zeitschrift für Medizinische Physik* **29**(2), 86–101 (2019). <https://doi.org/10.1016/j.zemedi.2018.12.003>
8. Maier, A., Zaiss, M., Bayer, S.: Beating the style detector: Three hours of agentic research on the AI-text arms race. arXiv preprint arXiv:2605.02620 (2026)
9. OWASP GenAI Security Project: OWASP top 10 for LLM applications: LLM01:2025 prompt injection (2025), <https://genai.owasp.org/llmrisk/llm01-prompt-injection/>
10. Salvi, F., Horta Ribeiro, M., Gallotti, R., West, R.: On the conversational persuasiveness of GPT-4. *Nature Human Behaviour* **9**(8), 1645–1653 (2025). <https://doi.org/10.1038/s41562-025-02194-6>
11. Semmelrock, H., Ross-Hellauer, T., Kopeinik, S., et al.: Reproducibility in machine-learning-based research: Overview, barriers, and drivers. *AI Magazine* **46**(2), e70002 (2025). <https://doi.org/10.1002/aaai.70002>
12. Wallace, E., Xiao, K., Leike, R., Weng, L., Heidecke, J., Beutel, A.: The instruction hierarchy: Training LLMs to prioritize privileged instructions. arXiv preprint arXiv:2404.13208 (2024)
13. Weissburg, I., Anand, S., Levy, S., Jeong, H.: LLMs are biased teachers: Evaluating LLM bias in personalized education. In: *Findings of the Association for Computational Linguistics: NAACL 2025* (2025), <https://aclanthology.org/2025.findings-naacl.314/>
14. Wu, A.J., Liu, R., Li, S.S., Tsvetkov, Y., Griffiths, T.L.: Ads in AI chatbots? an analysis of how large language models navigate conflicts of interest. arXiv preprint arXiv:2604.08525 (2026)