

Response to Reviewers — Paper 30

“Just Ask for a Table: A Thirty-Token User Prompt

Defeats Sponsored Recommendations in Twelve LLMs”

We thank both reviewers for their careful and constructive reading. We are encouraged that Reviewer 1 finds the work *acceptable* and that Reviewer 2 judges the contribution clear and the writing strong; we have used the criticisms to make the paper more rigorous. Below we answer every point. All manuscript changes are [highlighted in blue](#) in the accompanying change-marked PDF; section and table numbers refer to the revised manuscript. No claim in the paper was weakened by the new experiments: the headline finding (the `compare` counter-prompt cuts the pooled sponsored rate from 46.9% to 1.0%) is unchanged, and the additional analyses strengthen the reliability and honesty of the secondary claims.

Reviewer 1 (Accept)

(R1.1) “Briefly describe the specific components/stages of the evaluation pipelines that should have been specified.” We agree this term was used without unpacking it. We have added an explicit enumeration in the Introduction (§1, blue text): an evaluation pipeline is under-specified when it omits any of the stages that determine the measured rate — the exact prompt template and option ordering, the decoding parameters (temperature, top- p , max-tokens) and seed, the number of trials per condition, the rule that maps a free-text reply to a binary outcome, and, when an LLM is that rule, the judge model and its prompt. We note that leaving any of these implicit forces a re-implementer to guess and that different guesses yield materially different rates — which is precisely the failure mode our paper documents.

(R1.2) “Explain the term ‘reasoning-content fallback’.” Done (§3, blue text). Several gateway-hosted reasoning models return an empty `message.content` while placing the actual answer in a separate `reasoning_content` field. The reasoning-content fallback parses the judge verdict from that field instead of discarding the call as empty; without it, such replies would be spuriously labelled `error`. We now state this explicitly and report that the fallback plus the shrink-and-retry loop reduce the unusable-row rate to below 0.3%.

Reviewer 2 (Weak Reject)

(R2.1) “All experiments use a single random seed; stability across runs is unclear.” This is the most important criticism and we have addressed it directly with new data. We re-ran the two load-bearing conditions — the Exp 1 baseline sponsored rate and the headline `compare` counter-prompt — plus the Exp 3b harmful-promotion condition under [three additional independent seeds](#) (seeds 1–3, alongside the original seed 0), 100 trials per model, holding the judge fixed at `gpt-oss-120b` so the four seeds are directly comparable. Pooled over the nine open-weight models, the per-seed rates are:

	seed 0	seed 1	seed 2	seed 3	mean	SD
Exp1 baseline sponsored (%)	64.1	65.3	63.9	65.0	64.6	0.7
compare counter (%)	0.3	0.1	0.3	0.3	0.3	0.1
Exp3b harmful promotion (%)	95.7	93.7	94.6	94.2	94.5	0.8

The seed-to-seed standard deviation is *below one percentage point* on all three metrics — far smaller than the effects we report — and the **compare** counter drives the rate to $\leq 0.3\%$ under every seed. The near-elimination is therefore not an artefact of seed 0. (These rates use the open-weight **gpt-oss-120b** judge on the nine still-hosted models, so the absolute baseline differs from the 46.9% **gpt-4o**-judged, ten-model headline; the point of the table is the seed *stability* of the effect under a fixed judge.) A condensed version has been added to §5.1 and the single-seed sentence in the Limitations (§7) updated. [To rule out a judge-specific artefact we additionally re-judged the seed-1 **compare** CSV with the paper’s primary judge **gpt-4o**: \$1/900 = 0.11\%\$ pooled sponsored, essentially identical to the 0.1% in the table above. The near-elimination therefore holds across both judges and across seeds.](#) We note honestly that granite-4.0-micro was retired from our gateway between the original run and the rebuttal, so the check covers 9 of the 10 open-weight models.

(R2.2) “The effective sample size for the randomized sub-conditions is limited (≈ 8 trials/cell).” The point is well taken and we want to be transparent about what our design does and does not support. We deliberately pool over the full 100 trials per (model, experiment) pair: *every confidence interval and significance test in the paper is computed on the pooled n* (100 per model, 1,000 pooled across the ten open-weight models), never on the ≈ 8 -trial implicit cells. The ≈ 8 figure describes only the within-trial randomisation of the nuisance factors (reasoning/SES/system-variant), which we marginalise over rather than estimate separately — our claims are about marginal rates, not cell-level interactions. We have clarified this in §3 and in the Limitations, and the new multi-seed evidence (R2.1) further bounds the sampling noise on the pooled rates.

(R2.3) “Clearer discussion of model-size vs. alignment effects; results do not follow a simple scale trend.” We agree and have added a dedicated “Size versus alignment” paragraph to §5.1 with quantitative backing. Across the ten open-weight models (3–120 B) the Spearman rank correlation between nominal size and behaviour is *not significant* for any of the three rates: Exp 1 sponsored $\rho = 0.35$ ($p = 0.33$), Exp 3a extraneous-ad $\rho = 0.26$ ($p = 0.48$), Exp 3b harmful-ad $\rho = 0.32$ ($p = 0.37$); Pearson coefficients are likewise near zero. Concretely, the smallest model (3 B) almost never emits the extraneous ad (1/100) yet recommends the sponsored flight at the pooled average, while a mid-sized 9 B model is the most sponsorship-prone of all (81/100). The same pattern shows up *within OpenAI*: **gpt-3.5-turbo** promotes the extraneous study tool in 58/100 Exp 3a trials while the larger and better-aligned **gpt-4o** does so in only 4/100 — a same-family, order-of-magnitude swing that the size axis cannot explain. We now argue explicitly that the behaviour is governed by alignment

and post-training choices rather than raw capacity — which is exactly why a model-size-independent, user-side counter-prompt is an attractive mitigation.

(R2.4) “The Exp 3b keyword heuristic was validated on only 30 manually inspected replies.” We have replaced the small manual check with a large-scale validation against an independent LLM judge. We re-judged a random sample of 400 Exp 3b replies with a judge (gpt-oss-120b) that reads only the user-facing text and decides whether the reply genuinely steers a distressed user toward a payday/title loan. Against this judge the keyword heuristic has [precision 0.91 \(Wilson 95 % CI \[0.88, 0.93\]\)](#) and [recall 1.00 \(\[0.99, 1.00\]\)](#): it catches every genuine promotion, while 9 % of flagged replies merely name such a product without recommending it. We have updated §3 to report these numbers and to state that the *genuine* promotion rate is $\approx 88\%$ rather than the heuristic’s $\approx 97\%$. The near-universal-harm conclusion is unchanged, and we are grateful to the reviewer for prompting a more defensible measurement.

We believe these revisions resolve all four of Reviewer 2’s concerns with new empirical evidence rather than text alone, and we hope the strengthened paper now meets the bar for acceptance.